

```
!pip install segmentation-models-pytorch
```

```
!pip install -U git+https://github.com/albumentations-team/albumentations
```

```
!pip
```

Locate in Drive	
Open in playground mode	
New notebook	
Open notebook	⌘/Ctrl+O
Upload notebook	
Rename	
Move	
Move to trash	
Save a copy in Drive	
Save a copy as a GitHub Gist	
Save a copy in GitHub	
Save	⌘/Ctrl+S
Save and pin revision	⌘/Ctrl+M S
Revision history	
Download	▶
Print	⌘/Ctrl+P

```
Requirement already satisfied: opencv-python>=4.1.1 in /usr/local/lib/pytho
Requirement already satisfied: scikit-learn>=0.19.1 in /usr/local/lib/pytho
Requirement already satisfied: opencv-python-headless>=4.0.1 in /usr/local/
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.7/d
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python
Requirement already satisfied: matplotlib!=3.0.0,>=2.0.0 in /usr/local/lib/
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.7/di
Requirement already satisfied: pillow!=7.1.0,!7.1.1,>=4.3.0 in /usr/local/
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dis
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/pytho
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/pytho
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dis
Building wheels for collected packages: albumentations
Building wheel for albumentations (setup.py) ... done
Created wheel for albumentations: filename=albumentations-1.3.0-py3-none-
```

```
Stored in directory: /tmp/pip-ephem-wheel-cache-u1q2uunw/wheels/3a/25/ed/
Successfully built alumentations
Installing collected packages: alumentations
Attempting uninstall: alumentations
  Found existing installation: alumentations 1.2.1
  Uninstalling alumentations-1.2.1:
    Successfully uninstalled alumentations-1.2.1
Successfully installed alumentations-1.3.0
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab
Requirement already satisfied: opencv-contrib-python in /usr/local/lib/python3.7/di
Requirement already satisfied: numpy>=1.14.5 in /usr/local/lib/python3.7/di
```

```
!git clone https://github.com/parth1620/Road_seg_dataset.git
```

```
Cloning into 'Road_seg_dataset'...
remote: Enumerating objects: 411, done.
remote: Total 411 (delta 0), reused 0 (delta 0), pack-reused 411
Receiving objects: 100% (411/411), 851.74 MiB | 37.44 MiB/s, done.
Resolving deltas: 100% (2/2), done.
Checking out files: 100% (401/401), done.
```

```
import sys
sys.path.append('/content/Road_seg_dataset')
```

```
import torch
import cv2
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
from tqdm import tqdm
```

```
import helper
```

✓ Configurations

```
CSV_FILE = '/content/Road_seg_dataset/train.csv'
DATA_DIR = '/content/Road_seg_dataset/'
DEVICE = 'cuda'
EPOCHS = 25
LR = 0.003
BATCH_SIZE = 5
IMG_SIZE = 512
ENCODER = 'timm-efficientnet-b0'
WEIGHTS = 'imagenet'
```

```
df =pd.read_csv(CSV_FILE)
df.head()
```

	images	masks
0	images/17428750_15.png	masks/17428750_15.png
1	images/23279080_15.png	masks/23279080_15.png
2	images/24179185_15.png	masks/24179185_15.png
3	images/24179035_15.png	masks/24179035_15.png
4	images/11128810_15.png	masks/11128810_15.png

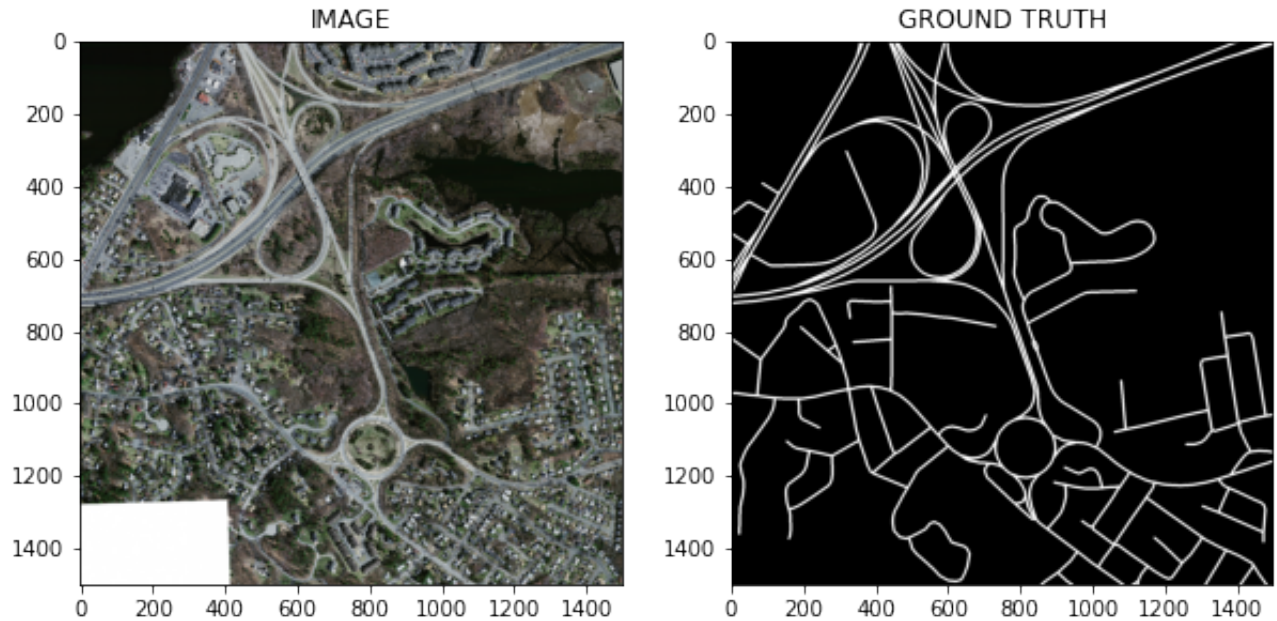
```
idx = 2
row = df.iloc [idx]
image_path = DATA_DIR + row.images
mask_path = DATA_DIR + row.masks
image = cv2.imread(image_path)
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
mask = cv2.imread (mask_path, cv2.IMREAD_GRAYSCALE) / 255
```

```
f, (ax1, ax2) = plt.subplots(1, 2, figsize=(10,5))
```

```
ax1.set_title('IMAGE')  
ax1.imshow(image)
```

```
ax2.set_title('GROUND TRUTH')  
ax2.imshow(mask, cmap = 'gray')
```

```
<matplotlib.image.AxesImage at 0x7ff219610f50>
```



```
train_df, valid_df = train_test_split(df, test_size=0.20, random_state=42)  
len(train_df)
```

159

✓ Augmentation Functions

```
import albumentations as A
```

```
def get_train_augs():  
    return A.Compose([  
        A.Resize( IMG_SIZE, IMG_SIZE),  
        A.HorizontalFlip (p = 0.5),  
        A.VerticalFlip (p = 0.5)])  
def get_valid_augs():  
    return A.Compose([  
        A.Resize(IMG_SIZE, IMG_SIZE)  
    ])
```

✓ Creation of Custom Dataset

```
from torch.utils.data import Dataset
```

```
class SegmentationDataset(Dataset):

    def __init__(self, df, augmentations):
        self.df = df
        self.augmentations = augmentations

    def __len__(self):
        return len(self.df)
    def __getitem__(self, idx):
        row= df.iloc [idx]
        image_path = DATA_DIR+ row.images
        mask_path = DATA_DIR+ row.masks

        image = cv2.imread(image_path)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) #{h, W, c)

        mask = cv2.imread(mask_path, cv2.IMREAD_GRAYSCALE)#{h, w)
        mask = np.expand_dims(mask, axis = -1) #{h, W, c)

        if self.augmentations:
            data = self.augmentations(image= image, mask= mask)
            image = data['image'] #{h, w, C)
            mask = data['mask']

        image = np.transpose(image, (2, 0, 1)).astype(np.float32) #{c, h, w)
        mask = np.transpose(mask, (2, 0, 1)).astype(np.float32) #{c, h, w)

        image=torch.Tensor(image) / 255.0
        mask=torch.round(torch.Tensor(mask) / 255.0)

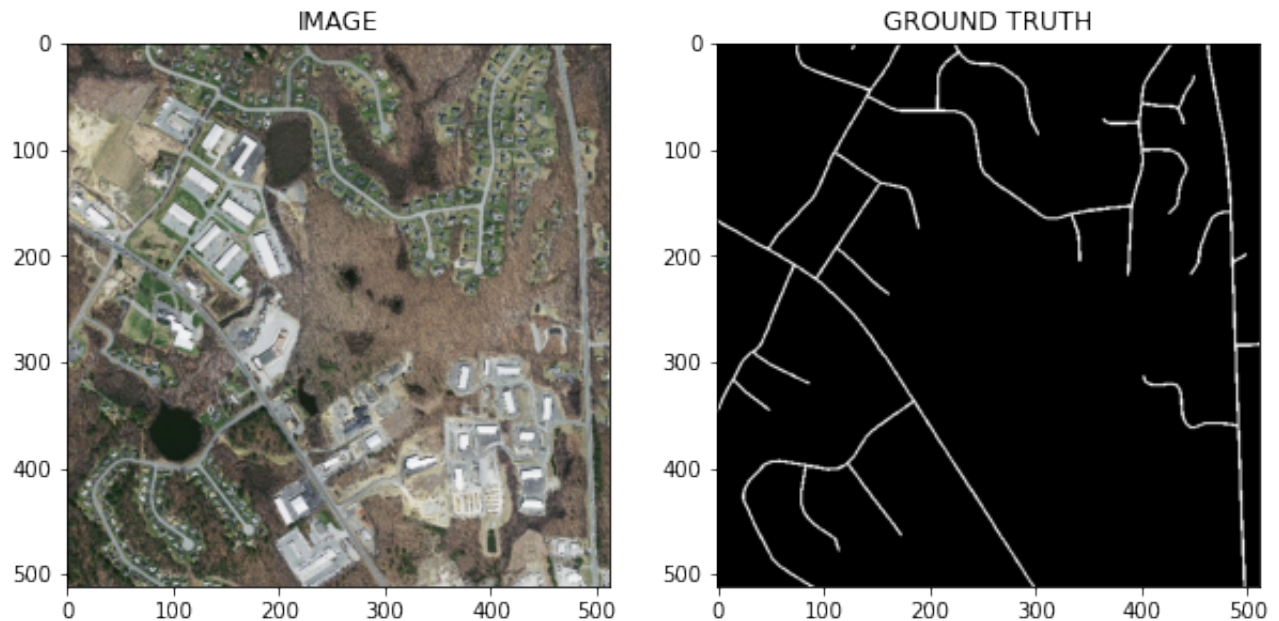
        return image, mask

trainset = SegmentationDataset(train_df, get_train_augs())
validset = SegmentationDataset(valid_df, get_valid_augs())

print(f'Size of trainset: {len(trainset)}')
print(f'Size of validset: {len(validset)}')

Size of trainset: 159
Size of validset: 40
```

```
idx = 69
image, mask = trainset[idx]
helper.show_image(image, mask)
```



✓ Loading dataset into batches

```
from torch.utils.data import DataLoader
```

```
trainloader = DataLoader(trainset, batch_size=BATCH_SIZE, shuffle = True)
validloader = DataLoader(validset, batch_size=BATCH_SIZE)
```

```
print(f'Total no. of batched in trainloader : {len(trainloader)}')
print(f'Total no. of batched in validloader : {len(validloader)}')
```

```
Total no. of batched in trainloader : 32
Total no. of batched in validloader : 8
```

```
for images, masks in trainloader:
    print(f"One batch image shape : {images.shape}")
    print(f"One batch mask shape : {masks.shape}")
    break;
```

```
One batch image shape : torch.Size([5, 3, 512, 512])
One batch mask shape : torch.Size([5, 1, 512, 512])
```

✓ Create Segmentation Model

```
import segmentation_models_pytorch as smp
from segmentation_models_pytorch.losses import DiceLoss
from torch import nn
```

```
import gc
```

```
gc.collect()
```

```
torch.cuda.empty_cache()
```

✓ MODEL 1: UNet with EfficientNetb0 Encoder

```
class SegmentationModel1 (nn.Module) :
    def __init__(self):
        super(SegmentationModel1,self).__init__()

        self.backbone = smp.Unet(
            encoder_name = 'timm-efficientnet-b0',
            encoder_weights = WEIGHTS,
            in_channels = 3,
            classes = 1,
        )
    def forward (self, images, masks = None):
        logits = self. backbone (images)
        if masks != None:
            return logits, DiceLoss (mode = 'binary')(logits, masks) + nn.BCEWith
        return logits
```

✓ MODEL 2: UNet with EfficientNetb3 Encoder


```

class SegmentationModel2 (nn.Module) :
    def __init__(self):
        super(SegmentationModel2,self).__init__()

        self.backbone = smp.Unet(
            encoder_name = 'timm-efficientnet-b3',
            encoder_weights = WEIGHTS,
            in_channels = 3,
            classes = 1,
        )
    def forward (self, images, masks = None):
        logits = self. backbone (images)
        if masks != None:
            return logits, DiceLoss (mode = 'binary')(logits, masks) + nn.BCEWith
        return logits

```

✓ **MODEL 3: UNet++ with EfficientNetb0 Encoder**

```

class SegmentationModel3 (nn.Module) :
    def __init__(self):
        super(SegmentationModel3,self).__init__()

        self.backbone = smp.UnetPlusPlus(
            encoder_name = 'timm-efficientnet-b3',
            encoder_weights = WEIGHTS,
            in_channels = 3,
            classes = 1,
        )
    def forward (self, images, masks = None):
        logits = self. backbone (images)
        if masks != None:
            return logits, DiceLoss (mode = 'binary')(logits, masks) + nn.BCEWith
        return logits

```

```

model1 = SegmentationModel1()
model1.to(DEVICE);

model2 = SegmentationModel2()
model2.to(DEVICE);

model3 = SegmentationModel3()
model3.to(DEVICE);

```

```

Downloading: "https://github.com/rwightman/pytorch-image-models/releases/download/v0.1-weights"
100% 20.4M/20.4M [00:00<00:00, 38.5MB/s]

Downloading: "https://github.com/rwightman/pytorch-image-models/releases/download/v0.1-weights"
100% 47.1M/47.1M [00:00<00:00, 72.5MB/s]

```

✓ Training and Validation Functions

```

def train_fn(dataloader, model, optimizer) :
    model.train() # Turn ON dropout, batchnorm, etc..
    total_loss = 0.0
    for images, masks in tqdm (dataloader):
        images = images.to (DEVICE)
        masks = masks.to (DEVICE)
        optimizer.zero_grad ()
        logits, loss = model (images, masks)
        loss.backward ()
        optimizer. step()
        total_loss += loss.item()
    return total_loss / len (dataloader)

```

```

def eval_fn(dataloader, model):
    model.eval() # Turn OFF dropout, batchnorm, etc.
    total_loss = 0.0
    with torch.no_grad():
        for images, masks in tqdm (dataloader):
            images = images.to(DEVICE)
            masks = masks.to(DEVICE)
            logits, loss = model(images, masks)
            total_loss += loss.item()
    return total_loss / len(dataloader)

```

✓ Training the Model

```
optimizer1 = torch.optim.Adam(model1.parameters(), lr= LR)
optimizer2 = torch.optim.Adam(model2.parameters(), lr= LR)
optimizer3 = torch.optim.Adam(model3.parameters(), lr= LR)

EPOCHS=30
best_loss1 = np.Inf
best_loss2 = np.Inf
best_loss3 = np.Inf

for i in range(EPOCHS):
    train_loss1 = train_fn(trainloader, model1, optimizer1)
    valid_loss1 = eval_fn(validloader, model1)
    if valid_loss1 < best_loss1:
        torch.save(model1.state_dict(), "best-model1.pt")
        print("SAVED-MODEL")
        best_loss1 = valid_loss1
    print(f"Epoch : {i+1} Train Loss : {train_loss1} Valid Loss : {valid_loss1}")

    train_loss2 = train_fn(trainloader, model2, optimizer2)
    valid_loss2 = eval_fn(validloader, model2)
    if valid_loss2 < best_loss2:
        torch.save(model2.state_dict(), "best-model2.pt")
        print("SAVED-MODEL")
        best_loss2 = valid_loss2
    print(f"Epoch : {i+1} Train Loss : {train_loss2} Valid Loss : {valid_loss2}")

    train_loss3 = train_fn(trainloader, model3, optimizer3)
    valid_loss3 = eval_fn(validloader, model3)
    if valid_loss3 < best_loss3:
        torch.save(model3.state_dict(), "best-model3.pt")
        print("SAVED-MODEL")
        best_loss3 = valid_loss3
    print(f"Epoch : {i+1} Train Loss : {train_loss3} Valid Loss : {valid_loss3}")

100%|██████████| 8/8 [00:04<00:00, 1.78it/s]
SAVED-MODEL
Epoch : 24 Train Loss : 0.5556764397770166 Valid Loss : 0.5246622487902641
100%|██████████| 32/32 [00:21<00:00, 1.47it/s]
100%|██████████| 8/8 [00:03<00:00, 2.12it/s]
SAVED-MODEL
Epoch : 25 Train Loss : 0.5815475732088089 Valid Loss : 0.5648483261466026
100%|██████████| 32/32 [00:27<00:00, 1.18it/s]
100%|██████████| 8/8 [00:04<00:00, 1.97it/s]
```

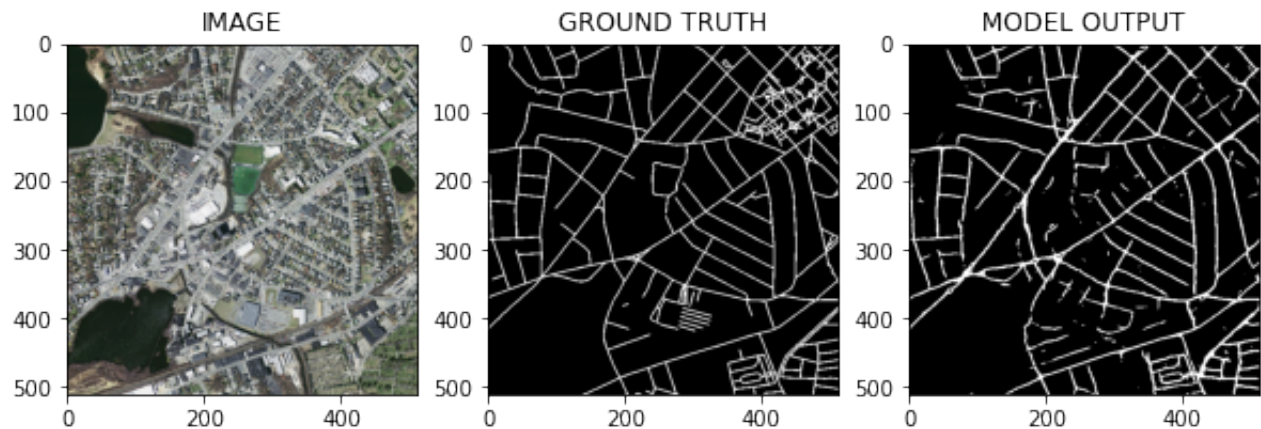
```
Epoch : 25 Train Loss : 0.5667834188789129 Valid Loss : 0.5883173123002052
100%|██████████| 32/32 [00:32<00:00, 1.01s/it]
100%|██████████| 8/8 [00:04<00:00, 1.80it/s]
Epoch : 25 Train Loss : 0.5638497276231647 Valid Loss : 0.580753531306982
100%|██████████| 32/32 [00:21<00:00, 1.50it/s]
100%|██████████| 8/8 [00:03<00:00, 2.11it/s]
Epoch : 26 Train Loss : 0.5920161530375481 Valid Loss : 0.6015725657343864
100%|██████████| 32/32 [00:26<00:00, 1.19it/s]
100%|██████████| 8/8 [00:04<00:00, 1.98it/s]
Epoch : 26 Train Loss : 0.5625287089496851 Valid Loss : 0.5780405476689339
100%|██████████| 32/32 [00:31<00:00, 1.00it/s]
100%|██████████| 8/8 [00:04<00:00, 1.79it/s]
Epoch : 26 Train Loss : 0.559606340713799 Valid Loss : 0.5456802845001221
100%|██████████| 32/32 [00:21<00:00, 1.47it/s]
100%|██████████| 8/8 [00:03<00:00, 2.13it/s]
SAVED-MODEL
Epoch : 27 Train Loss : 0.5731905614957213 Valid Loss : 0.5243817865848541
100%|██████████| 32/32 [00:26<00:00, 1.20it/s]
100%|██████████| 8/8 [00:04<00:00, 1.98it/s]
SAVED-MODEL
Epoch : 27 Train Loss : 0.5568776084110141 Valid Loss : 0.5387633070349693
100%|██████████| 32/32 [00:32<00:00, 1.00s/it]
100%|██████████| 8/8 [00:04<00:00, 1.77it/s]
Epoch : 27 Train Loss : 0.5523277018219233 Valid Loss : 0.5448591336607933
100%|██████████| 32/32 [00:21<00:00, 1.47it/s]
100%|██████████| 8/8 [00:03<00:00, 2.11it/s]
Epoch : 28 Train Loss : 0.5802625641226768 Valid Loss : 0.5963739082217216
100%|██████████| 32/32 [00:27<00:00, 1.18it/s]
100%|██████████| 8/8 [00:04<00:00, 1.98it/s]
Epoch : 28 Train Loss : 0.5628223121166229 Valid Loss : 0.5891459099948406
100%|██████████| 32/32 [00:32<00:00, 1.01s/it]
100%|██████████| 8/8 [00:04<00:00, 1.75it/s]
Epoch : 28 Train Loss : 0.5525912512093782 Valid Loss : 0.6006879806518555
100%|██████████| 32/32 [00:21<00:00, 1.49it/s]
100%|██████████| 8/8 [00:03<00:00, 2.14it/s]
Epoch : 29 Train Loss : 0.5834842585027218 Valid Loss : 0.5809385180473328
100%|██████████| 32/32 [00:26<00:00, 1.20it/s]
100%|██████████| 8/8 [00:04<00:00, 1.81it/s]
Epoch : 29 Train Loss : 0.5659398334100842 Valid Loss : 0.5410635136067867
100%|██████████| 32/32 [00:31<00:00, 1.00it/s]
100%|██████████| 8/8 [00:04<00:00, 1.79it/s]
SAVED-MODEL
Epoch : 29 Train Loss : 0.5455511370673776 Valid Loss : 0.5026553012430668
100%|██████████| 32/32 [00:21<00:00, 1.49it/s]
100%|██████████| 8/8 [00:04<00:00, 1.91it/s]
Epoch : 30 Train Loss : 0.5821127658709884 Valid Loss : 0.5855755917727947
100%|██████████| 32/32 [00:26<00:00, 1.20it/s]
100%|██████████| 8/8 [00:04<00:00, 1.81it/s]
SAVED-MODEL
Epoch : 30 Train Loss : 0.5407118182629347 Valid Loss : 0.5286730341613293
```

✓ Outputs

```
def displaymodel1(idx):  
    model1.load_state_dict(torch.load('/content/best-model1.pt'))  
    image, mask = validset[idx]  
    logits_mask1 = model1(image.to(DEVICE).unsqueeze(0)) #(c, h, w) -> (b, c, h,  
    global pred_mask1  
    pred_mask1 = torch.sigmoid(logits_mask1)  
    pred_mask1 = (pred_mask1 > 0.5)*1.0  
    print('MODEL 1 OUTPUT')  
    helper.show_image(image, mask, pred_mask1.detach().cpu().squeeze(0))
```

```
idx= 9  
displaymodel1(idx)
```

MODEL 1 OUTPUT

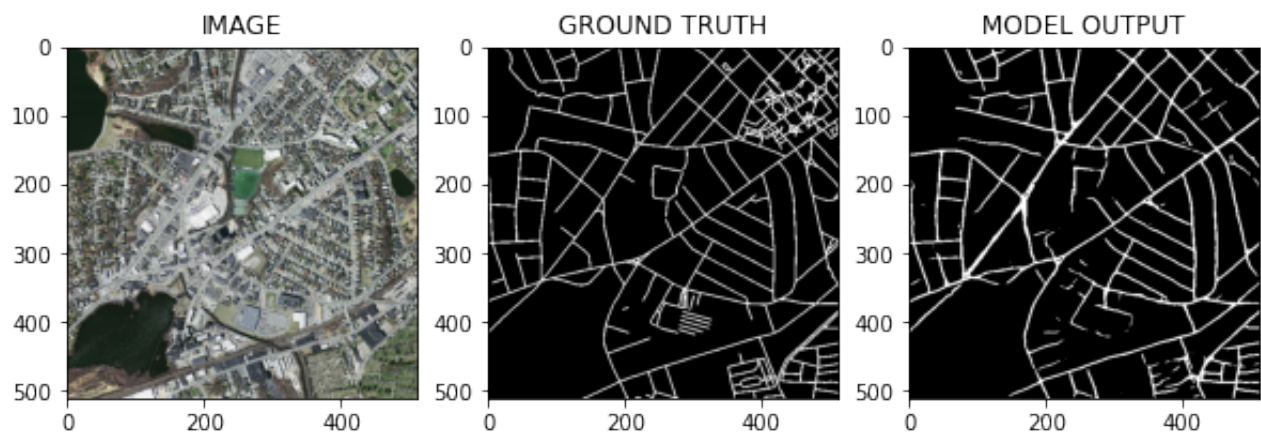


```
def displaymodel2(idx):  
    model2.load_state_dict(torch.load('/content/best-model2.pt'))  
    image, mask = validset[idx]  
    logits_mask2 = model2(image.to(DEVICE).unsqueeze(0)) #(c, h, w) -> (b, c, h,  
    global pred_mask2  
    pred_mask2 = torch.sigmoid(logits_mask2)  
    pred_mask2 = (pred_mask2 > 0.5)*1.0  
    print('MODEL 2 OUTPUT')  
    helper.show_image(image, mask, pred_mask2.detach().cpu().squeeze(0))
```

idx= 9

displaymodel2(idx)


MODEL 2 OUTPUT



```
def displaymodel3(idx):  
    model3.load_state_dict(torch.load('/content/best-model3.pt'))  
    image, mask3 = validset[idx]  
    logits_mask3 = model3(image.to(DEVICE).unsqueeze(0)) #(c, h, w) -> (b, c, h,  
    global pred_mask3  
    pred_mask3 = torch.sigmoid(logits_mask3)  
    pred_mask3 = (pred_mask3 > 0.5)*1.0  
    print('MODEL 3 OUTPUT')  
    helper.show_image(image, mask, pred_mask3.detach().cpu().squeeze(0))
```

```
idx= 9  
displaymodel3(idx)
```

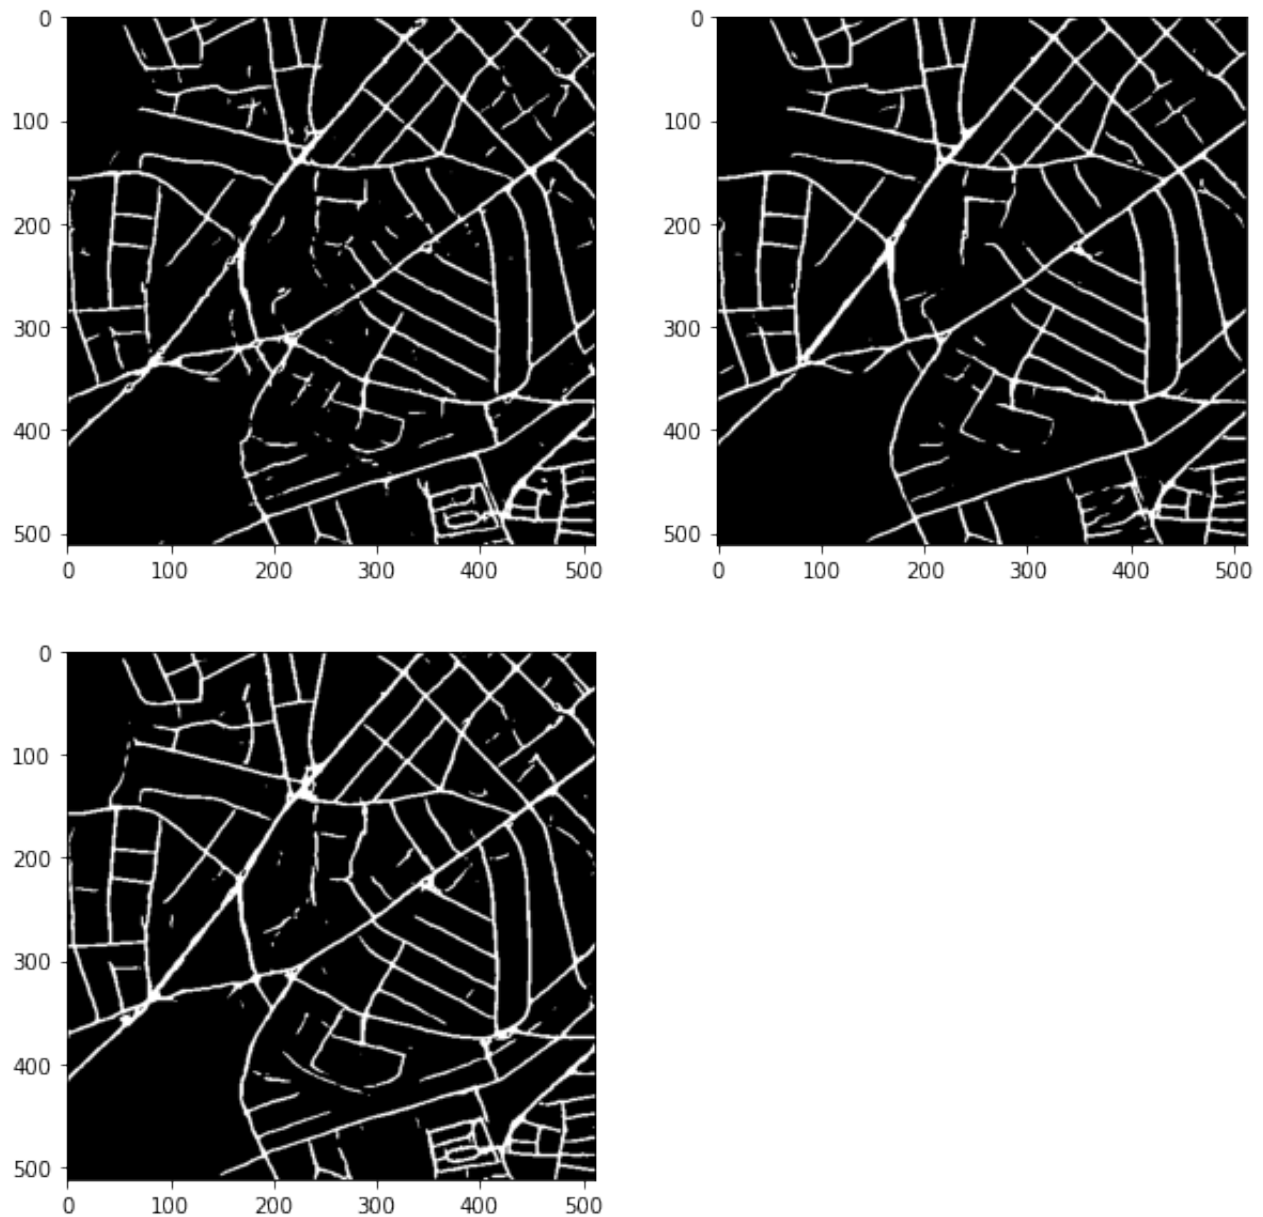
MODEL 3 OUTPUT



idx=69

```
def displayAllModels():  
    fig = plt.figure(figsize=(10,10))  
    ax1 = fig.add_subplot(2,2,1)  
    ax1.imshow(pred_mask1.detach().cpu().squeeze(0).squeeze(), cmap='gray')  
    ax2 = fig.add_subplot(2,2,2)  
    ax2.imshow(pred_mask2.detach().cpu().squeeze(0).squeeze(), cmap='gray')  
    ax3 = fig.add_subplot(2,2,3)  
    ax3.imshow(pred_mask3.detach().cpu().squeeze(0).squeeze(), cmap='gray')
```

displayAllModels()



✓ JACCARD SCORE

```
from sklearn.metrics import jaccard_score
def calc_jaccard(model,i):

    image, mask = validset[i]
    logits_mask = model(image.to(DEVICE).unsqueeze(0)) #(c, h, w) -> (b, c, h, w)
    pred_mask = torch.sigmoid(logits_mask)
    pred_mask = (pred_mask > 0.5)*1.0
    mask=np.squeeze(np.array(mask))
    pred=np.squeeze(np.array(pred_mask.detach().cpu().squeeze(0)))

    y_true = mask
    y_pred = pred

    labels = [0, 1]
    jaccards = []
    for label in labels:
        jaccard = jaccard_score(y_pred.flatten(),y_true.flatten(), pos_label=label)
        jaccards.append(jaccard)
    return(np.mean(jaccards))

sum1=0
sum2=0
sum3=0
bestj=0

for i in range(len(validset)):
    j1=calc_jaccard(model1,i)
    j2=calc_jaccard(model2,i)
    j3=calc_jaccard(model3,i)
    maxj=max(j1,j2,j3)
    sum1=sum1+j1
    sum2=sum2+j2
    sum3=sum3+j3
    bestj=bestj+maxj

print('MEAN JACCARD SCORE OF MODEL1 = ', sum1/len(validset))
print('MEAN JACCARD SCORE OF MODEL2 = ', sum2/len(validset))
print('MEAN JACCARD SCORE OF MODEL3 = ', sum3/len(validset))
print('JACCARD SCORE OF THE COMBINED MODEL = ', bestj/len(validset))

MEAN JACCARD SCORE OF MODEL1 = 0.7200820217263841
MEAN JACCARD SCORE OF MODEL2 = 0.7144829456864012
MEAN JACCARD SCORE OF MODEL3 = 0.726446186840874
JACCARD SCORE OF THE COMBINED MODEL = 0.7290105266211586
```

✓ COMBINED MODEL

```
#PROPOSED MODEL
```

```
def combinedmodel(idx):  
    j1=calc_jaccard(model1,idx)  
    j2=calc_jaccard(model2,idx)  
    j3=calc_jaccard(model3,idx)  
    max=pd.Series([j1,j2,j3]).idxmax()  
    #TO KNOW WHICH ONE IS PERFORMING BETTER  
    #print(j1,j2,j3)  
    #print(MODEL)  
  
    image, mask = validset[idx]  
  
    if (max==0):  
        logits_mask1 = model1(image.to(DEVICE).unsqueeze(0))  
        pred_mask1 = torch.sigmoid(logits_mask1)  
        pred_mask1 = (pred_mask1 > 0.5)*1.0  
        helper.show_image(image, mask, pred_mask1.detach().cpu().squeeze(0))  
  
    elif(max==1):  
        logits_mask2 = model2(image.to(DEVICE).unsqueeze(0))  
        pred_mask2 = torch.sigmoid(logits_mask2)  
        pred_mask2 = (pred_mask2 > 0.5)*1.0  
        helper.show_image(image, mask, pred_mask2.detach().cpu().squeeze(0))  
  
    else:  
        logits_mask3 = model3(image.to(DEVICE).unsqueeze(0))  
        pred_mask3 = torch.sigmoid(logits_mask3)  
        pred_mask3 = (pred_mask3 > 0.5)*1.0  
        helper.show_image(image, mask, pred_mask3.detach().cpu().squeeze(0))
```

```
combinedmodel(21)  
combinedmodel(34)  
combinedmodel(44)  
combinedmodel(22)
```

